WEARIFY

AI VIRTUAL TRY-ON PLATFORM · INDIAN SAREE RETAIL

The Platform, the Move to Azure, and What It Costs

A complete walkthrough — the product, the architecture, the AI, the migration, and the budget.

OUR JOURNEY TODAY

What We'll Cover — Five Parts

1. **What Wearify is**
   The product and the six kinds of people who use it
2. **How it's built**
   The architecture, the technologies, and how they work together
3. **The AI try-on**
   How the "magic" actually works
4. **Moving to Microsoft Azure**
   Why, how, and the migration plan
5. **What it costs**
   The budget, the per-use cost, and our Azure credits

PART 1 · THE PRODUCT

What is Wearify?

An AI virtual try-on platform for Indian saree retailers.

- **The problem:** trying on saree after saree in a shop is slow and tiring.
- **The solution:** a customer stands at a smart mirror, picks any saree, and instantly sees a photorealistic image of themselves wearing it.
- **The scope:** built for saree stores across India — one platform serving six different kinds of users.

PART 1 · THE PRODUCT

One Platform, Six Users

Six different apps — all built from one codebase.

Admin	**Store**	**Tablet**	**Kiosk**	**Customer**	**Tailor**
Mission Control	**Retailer Dashboard**	**Sales Tablet**	**Smart Mirror**	**Mobile App**	**Tailor Portal**
Wearify operations	Shop owners	In-store staff	Customer, in-store	Customer, at home	Stitching partners

PART 2 · HOW IT'S BUILT

One App, One Backend

ONE Next.js application + ONE Convex backend

holds all the data — 46 tables

Admin | **Store** | **Tablet** | **Kiosk** | **Customer** | **Tailor**

Improve something shared — the data, a feature, a fix — and it reaches **all six surfaces at once**.

PART 2 · HOW IT'S BUILT

The Technologies We Use Today

TECHNOLOGY	WHAT IT IS	HOW WEARIFY USES IT
TypeScript & React	The programming language & UI library	Every screen and all the logic across the six surfaces is written in it
Next.js 16	A modern web framework	Builds and serves all six surfaces from a single codebase
Convex	A backend platform — database, server logic & real-time	Stores all 46 tables of data, runs the server functions, and pushes live updates
ComfyUI + Qwen models	An AI image-generation pipeline	Generates the virtual try-on image — the customer wearing the saree
RunPod	A serverless GPU cloud	Provides the graphics hardware (GPU) the AI pipeline runs on
TailwindCSS	A styling toolkit	Gives each of the six surfaces its distinct look and theme
@imgly background-removal	An in-browser AI tool	Cleans the background off the finished try-on image, on the device
SendGrid & Twilio	Email & messaging providers	Send email, SMS / OTP and WhatsApp messages to customers
Netlify	A web hosting platform	Serves the web app to users — fast, worldwide

PART 2 · HOW IT'S BUILT

How the Technologies Work Together

THE SIX SURFACES — ON USERS' DEVICES

Admin | **Store** | **Tablet** | **Kiosk** | **Customer** | **Tailor**

Next.js web app
the frontend — hosted today on Netlify

CONVEX BACKEND — the hub

Database	**File Storage**	**Real-time Engine**	**Scheduler**
46 tables of data	images & documents	pushes live updates	background jobs

RunPod GPU — ComfyUI + Qwen
runs the AI try-on and returns the image

SendGrid + Twilio
send email, SMS / OTP and WhatsApp

Convex is the **hub**: it stores the data, runs the logic, calls the AI and messaging services — and **pushes every change straight back to all six surfaces, live**.

PART 2 · HOW IT'S BUILT

The Big Idea: Live Data

The screen never has to ask "is it ready yet?"

✕ Most apps

The screen keeps **asking** the server — "any updates? any updates?" It feels slow and laggy.

✓ Wearify

The moment data changes, the server **pushes** it to every screen instantly — the mirror, the tablet and the phone all update together.

PART 3 · THE AI TRY-ON

What the AI Actually Does

Photo of the customer
a body scan at the mirror

+

Flat photo of a saree
a plain catalogue image

=

The customer wearing that saree
photorealistic — right border, colour & drape

PART 3 · THE AI TRY-ON

It's a Pipeline, Not One Model

A team of specialist AI models, passing work down a line — on a GPU.

1. Resize
prepare the two photos

2. Understand
Qwen-2.5-VL reads photos + instructions

3. Generate
Qwen-Image-Edit + LoRA paints the saree on

4. Decode
turn it into a normal photo

5. Clean up
remove background, in the browser

PART 3 · THE AI TRY-ON

One Try-On, End to End

1
Customer taps
"Try On" at the smart mirror.

2
App sends it
the request goes to the backend.

3
Backend checks
safety checks, then hands it to the GPU.

4
GPU generates
the AI pipeline runs and makes the image.

5
Live result
pushed back — the mirror updates by itself.

A few seconds, start to finish — everything in between is invisible to the customer.

PART 4 · MOVING TO MICROSOFT AZURE

Why Move to Azure?

Data stays in India

Azure's Central India region keeps customer data in-country — for the DPDP Act 2023.

One cloud

Replaces several separate providers with a single cloud, one bill.

Enterprise-grade

Proven security, scaling and reliability as Wearify grows.

Credits-funded

Microsoft startup credits cover the early journey (Part 5).

PART 4 · MOVING TO MICROSOFT AZURE

Two Ways to Migrate — and Our Choice

APPROACH A

Decompose

Break Convex into **six separate Azure services** and rebuild that part of the app. Powerful, but a re-architecture.

APPROACH B

Self-Host Convex

Keep Convex exactly as it is and **run it ourselves on Azure**. Nothing in the app is rewritten.

√ OUR CHOICE — lower risk, faster, no rewrite

PART 4 · MOVING TO MICROSOFT AZURE

The Migration Map

LAYER	TODAY	ON MICROSOFT AZURE
Web app hosting	Next.js on Netlify	**Azure App Service**
Backend & real-time	Convex Cloud	**Convex — self-hosted on Azure**
Database (under Convex)	managed by Convex	**Azure Database for PostgreSQL**
File storage	Convex File Storage	**Azure Blob Storage**
AI try-on (GPU)	ComfyUI + Qwen on RunPod	**Azure GPU — NVIDIA H100**
Sign-in	Better Auth / phone-OTP	**Microsoft Entra ID / External ID**
Email, SMS & WhatsApp	SendGrid + Twilio	**Azure Communication Services**
Region	US-based providers	**Central India (Pune)**

PART 4 · MOVING TO MICROSOFT AZURE

The AI on Azure — H100

The AI models do not change. The GPU steps up to the H100.

TODAY	→	ON AZURE
NVIDIA A100 (RunPod)	→	NVIDIA H100

Native FP8

Our models are FP8 — the H100 accelerates FP8 in hardware; the A100 cannot.

Faster renders

Higher throughput — shorter try-on times.

Room to scale

More concurrent try-ons per GPU as traffic grows.

PART 4 · MOVING TO MICROSOFT AZURE

How We'll Do It — Five Phases

PHASE 1

Azure foundation

Set up Azure in the Central India region.

PHASE 2

Self-hosted Convex

Stand up Convex on Azure + PostgreSQL.

PHASE 3

AI pipeline on H100

Move the try-on to Azure H100 GPUs.

PHASE 4

Frontend, sign-in & messaging

Move the web app, auth and messaging.

PHASE 5

Cutover

Run both side by side, switch, decommission.

Not one big switch — each phase is testable on its own.

PART 5 · WHAT IT COSTS

The Cost Model — Two Parts

Monthly cost = **Fixed baseline (~$450)** + (renders × **$0.05**) + messaging + storage

FIXED — always on

Keeps the platform switched on, even with zero customers.

~$450 / month

VARIABLE — scales with use

Grows with how many try-ons customers actually run.

~$0.05 / try-on

PART 5 · WHAT IT COSTS

Cost Per Use — the Unit Economics

$0.05

per AI try-on — about ₹4.
One customer visit (~4 try-ons) ≈ **$0.20 / ₹17**.
SMS, email, storage — rounding errors next to this.

Economies of scale: at low volume a render effectively costs 20¢+; as volume grows it drops toward the 5¢ floor. **The platform rewards usage.**

PART 5 · WHAT IT COSTS

Azure Credits — Burn Rate & Runway

Runway = Azure credits ÷ monthly burn. How long the credits last:

AZURE CREDITS	PILOT ~$725/MO	GROWTH ~$1,550/MO	SCALE ~$6,000/MO
$5,000	**~7 months**	**~3 months**	**under 1 month**
$25,000	**~3 years**	**~16 months**	**~4 months**
$150,000	**~17 years**	**~8 years**	**~2 years**

Credits also expire (~1 year after activation) — so we use them deliberately, not hoard them.

IN CLOSING

Five Things to Remember

1 **Wearify is an AI virtual try-on platform**
one codebase, six surfaces, one shared live database

2 **The AI is a pipeline of models on a GPU**
and it does not change when we move to Azure

3 **We keep Convex — self-hosted on Azure**
a relocation, not a rewrite

4 **H100 GPUs, in Central India**
faster AI, and DPDP Act 2023 compliant

5 **~$450 fixed + ~$0.05 per try-on**
and our Azure credits fund the journey

Thank you — questions welcome.

APPENDIX

Likely Questions

Q. How long will the migration take?
Five phases, each independently testable. A firm timeline comes after Phase 1 scoping; self-hosting Convex keeps it shorter than a full rebuild.

Q. Why H100 and not A100?
Our models are FP8, and the H100 accelerates FP8 in hardware — the A100 can't. Faster renders, and headroom to scale.

Q. Will there be downtime for customers?
No big-bang switch — old and new run side by side, and we cut over only when Azure is verified stable.

Q. What happens when the credits run out?
We plan to be revenue-generating or funded before then. The cost is predictable — ~$450 fixed + ~$0.05 per try-on.

Q. Will the AI quality change on Azure?
No. The same models and pipeline run on Azure — only the GPU changes. **Same input, same output.**

Q. Why keep Convex instead of native Azure services?
It avoids a rewrite. Self-hosting Convex is a relocation — lower risk, faster, and the live-data behaviour keeps working for free.

Q. How long do the Azure credits last?
Depends on usage — see slide 19. At pilot usage a $25,000 grant lasts ~3 years; at growth usage, ~16 months.

Q. Is the AI a third-party API, like ChatGPT?
No — we self-host open Qwen models in our own pipeline, so we control it fully. That's why it runs on our own GPUs.